

# NumPy

- **Introduces ndarrays.** Those are n-dimensional arrays of homogeneous data types, with many operations being performed on compiled code for performance.

## Working with NumPy

The first thing we need to do when installing a Python library we can type `pip install numpy` in the terminal. If we are using Jupyter Notebook we can run the following, which will basically call our terminal.

```
In [1]: !pip install numpy
```

Requirement already satisfied: numpy in c:\users\sergi\anaconda3\lib\site-packages (1.21.5)

Once NumPy is installed we might want to **import** it. To import a function we just need to type:

```
In [2]: import numpy as np
```

## Create an array from a list

```
In [4]: # Create a 1-dimensional array
```

```
A = np.array([1,3,4,5,6])    # A is our first array
```

```
print(A)  # allows to visualize A
type(A)   # check that the type is ndarray
```

```
[1 3 4 5 6]
numpy.ndarray
```

Out[4]:

```
In [6]: # Create a 2-dimensional array
```

```
B = np.array([[1,2],[3,4],[5,6]])    # 2-d array
print(B)
```

```
[[1 2]
 [3 4]
 [5 6]]
```

- **shape** : number of dimensions of the array.
- **size** : number of elements in the array
- **ndim** : number of dimensions in the array

```
In [11]: print(f'The dimensions of array A are{A.shape}')
print(f'The dimensions of array B are{B.shape}')
```

```
The dimensions of array A are(5,)
The dimensions of array B are(3, 2)
```

```
In [17]: # We can also check the number of dimnesions:

print(A.ndim)
print(B.ndim)

# And we can also check the type
print(A.dtype)

# We can also check the total number of elements of the array

print(A.size)
print(B.size)

1
2
int32
5
6
```

## Creating Arrays

### np.array()

Can take as inputs Python lists among others.

```
In [24]: list1 = [1,2,4]
list2 = [9,29129,12]
list3 = [list1,list2]

array1 = np.array(list3)
print(array1)
print(array1.dtype)

# if we wish we could specify the data type:

array2 = np.array(list3,dtype= np.int64)
print(array2.dtype)

[[ 1  2  4]
 [ 9 29129 12]]
int32
int64
```

### np.zeros()

It allows to create an array full of 0s.

```
In [26]: zeros = np.zeros(10)    # create a 10x1 array full of 0s.
print(zeros)

[0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
```

```
In [27]: zeros = np.zeros((10,3)) # create a 10x3 array full of 0s.
print(zeros)
```

```
[[0. 0. 0.]
 [0. 0. 0.]
 [0. 0. 0.]
 [0. 0. 0.]
 [0. 0. 0.]
 [0. 0. 0.]
 [0. 0. 0.]
 [0. 0. 0.]
 [0. 0. 0.]
 [0. 0. 0.]]
```

## np.ones()

It allows to create an array full of 1s.

```
In [28]: ones = np.ones(12)
print(ones)
```

```
[1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1.]
```

## np.empty()

It allows to create an array with random numbers. The reason to use empty over zeros is speed.

```
In [29]: empty = np.empty((2,2))
print(empty)
```

```
[[2.12199579e-314 4.67296746e-307]
 [5.98807563e-321 3.79442416e-321]]
```

## np.arange()

It allows to create arrays in a range of elements, considering the initial value, the final value, and the step size.

```
In [33]: C = np.arange(1,10,2)
print(C)

# A simpler version is
D = np.arange(4)
print(D)
```

```
[1 3 5 7 9]
[0 1 2 3]
```

## np.linspace()

It allows to create arrays with values that are spaced linearly in a specified interval. The difference with `np.arange()` is that here instead of step size we choose the number of steps.

```
In [34]: E = np.linspace(5,25,6)
print(E)
```

```
[ 5.  9. 13. 17. 21. 25.]
```

## np.ones\_like()

It allows to create an arrays of ones with the same shape as an specified array.

```
In [35]: array1 = np.array([[1,3],[3,4],[2,3]])

array2 = np.ones_like(array1)

print(array2)

[[1 1]
 [1 1]
 [1 1]]
```

The same can be done with `np.zeros_like()` and `np.empty_like()` .

## np.identity()

It allows to create the identity matrix.

```
In [38]: identity = np.identity(3)
print(identity)

[[1. 0. 0.]
 [0. 1. 0.]
 [0. 0. 1.]]
```

## Basic array operations

Operations are **element-wise** :

```
In [45]: a = np.array([20, 30, 40, 50])
b = np.arange(4)

#Sum:

print('The sum is :',a+b)
print('The sum is :',np.add(a,b))

# Subtraction

print('The subtraction is :',a-b)
print('The subtraction is :',np.subtract(a,b))
```

```
The sum is : [20 31 42 53]
The sum is : [20 31 42 53]
The subtraction is : [20 29 38 47]
The subtraction is : [20 29 38 47]
```

- **The product operator** `\*` is element-wise.
- **Matrix product** with `@`.

```
In [2]: # Multiplication:
A = np.array([[1, 1],
              [0, 1]])
B = np.array([[2, 0],
              [3, 4]])
```

```
# Element-wise:

print('element wise', A*B)
print(np.multiply(A,B))

# Matrix Product

print('matrix product', A@B)
print(A.dot(B))
```

```
element wise [[2 0]
 [0 4]]
[[2 0]
 [0 4]]
matrix product [[5 4]
 [3 4]]
[[5 4]
 [3 4]]
```

## Dvision

```
In [49]: # Dvision

# element-wise:
print(A/B)
print(np.divide(A,B))
```

```
[[0.5  inf]
 [0.   0.25]]
[[0.5  inf]
 [0.   0.25]]
```

```
C:\Users\Sergi\AppData\Local\Temp\ipykernel_23828\790273862.py:4: RuntimeWarning:
divide by zero encountered in true_divide
  print(A/B)
C:\Users\Sergi\AppData\Local\Temp\ipykernel_23828\790273862.py:5: RuntimeWarning:
divide by zero encountered in true_divide
  print(np.divide(A,B))
```

We can also perform operations between an array and a scalar.

```
In [51]: print(A**2)    # element-wise x^2
print(A*2)    # element-wise x*2
print(A+2)    # element-wise x+2
```

```
[[1 1]
 [0 1]]
[[2 2]
 [0 2]]
[[3 3]
 [2 3]]
```

## Indexing, Slicing and Iterating

Arrays can be indexed, sliced and iterated over much like lists and any other Python sequences. Let's see how it works with one-dimensional arrays:

```
In [55]: a = np.arange(10)**3
print(a)

print(a[2])    # access the element with index 2, the third element.
```

```
print(a[2:6]) # slice the array
print(a[-4:]) # slice from the -4 element to the end
```

```
[ 0  1  8 27 64 125 216 343 512 729]
8
[ 8 27 64 125]
[216 343 512 729]
```

We can also choose the step size in the interval slice. For example suppose we would like to obtain a subset containing only every three elements:

```
In [57]: print(a[1:8:3])
```

```
[ 1 64 343]
```

```
In [58]: # If we want to reverse the array:
print(a[::-1])
```

```
[729 512 343 216 125 64 27 8 1 0]
```

Multidimensional arrays have one index per axis. For example for a 2-dimensional array we would have `array[i][j]` where `i` are the row dimensions and `j` are the column dimensions. We can also write it as `array[i,j]`.

```
In [63]: a = np.array([[ 0, 1, 2, 3],
                      [10, 11, 12, 13],
                      [20, 21, 22, 23],
                      [30, 31, 32, 33],
                      [40, 41, 42, 43]])

# Access one element:

print(a[0,0],a[0][0]) # both will return the same

# Slices

print(a[:,2])          # only the third column
print(a[2:5,3:])

# We can also just include one index

print(a[2]) # third row
```

```
0 0
[ 2 12 22 32 42]
[[23]
 [33]
 [43]]
[20 21 22 23]
```

```
In [68]: # for a 3-dimensional array we have:
```

```
c = np.array([[[ 0, 1, 2],
               [ 10, 12, 13]],
              [[100, 101, 102],
               [110, 112, 113]]])

# we can access all its dimensions:

print(c[1,1,1])
print(c[1,...])
```

```
112
[[100 101 102]
 [110 112 113]]
```

The dots (...) represent as many colons as needed to produce a complete indexing tuple. For example, if `x` is an array with 5 axes, then `x[1, 2, ...]` is equivalent to `x[1, 2, :, :, :]`, `x[..., 3]` to `x[:, :, :, :, 3]` and `x[4, ..., 5, :]` to `x[4, :, :, 5, :]`.

## Slice based on logical conditions

We can also slice arrays based on logical conditions like `<` or `>`.

```
In [4]: a = np.array([[ 0,  1,  2,  3],
                    [10, 11, 12, 13],
                    [20, 21, 22, 23],
                    [30, 31, 32, 33],
                    [40, 41, 42, 43]])

print(a[a>10])    # only elements greater than 10
print(a[a<5])     # only elements smaller than 5
print(a[(a > 2) & (a < 11)]) # elements between 2 and 11

[11 12 13 20 21 22 23 30 31 32 33 40 41 42 43]
[0 1 2 3]
[ 3 10]
```

## Shape Manipulation and other methods

There are many ways of doing shape manipulation with arrays. We will now review some of them:

### `np.ravel()`

It returns the specified array flattened (with only one dimension).

```
In [5]: print(a)
print(a.ravel())    # the array is now flattened

[[ 0  1  2  3]
 [10 11 12 13]
 [20 21 22 23]
 [30 31 32 33]
 [40 41 42 43]]
[ 0  1  2  3 10 11 12 13 20 21 22 23 30 31 32 33 40 41 42 43]
```

### `np.reshape()`

The `reshape` method returns its argument with a modified shape.

```
In [84]: print(a.size) # the amount of elements needs to be the same
print(a)
print('New a ',a.reshape(2,5,2))
```

```

20
[[ 0  1  2  3]
 [10 11 12 13]
 [20 21 22 23]
 [30 31 32 33]
 [40 41 42 43]]
New a [[[ 0  1]
 [ 2  3]
 [10 11]
 [12 13]
 [20 21]]

 [[22 23]
 [30 31]
 [32 33]
 [40 41]
 [42 43]]]

```

## np.resize()

It does the same as `np.reshape` but instead of returning a new array, it modifies the array itself

```

In [85]: a.resize(5,2,2)
print(a)  # a is now modified.

```

```

[[[ 0  1]
 [ 2  3]]

 [[10 11]
 [12 13]]

 [[20 21]
 [22 23]]

 [[30 31]
 [32 33]]

 [[40 41]
 [42 43]]]

```

## np.swapaxes()

It allows to change two axis of an array.

```

In [86]: print(a)
b = np.swapaxes(a,0,1) # we are changing axis 0 for 1
print(b)

```



```
[[[ 0  1]
   [ 2  3]]
```

```
[[10 11]
 [12 13]]
```

```
[[20 21]
 [22 23]]
```

```
[[30 31]
 [32 33]]
```

```
[[40 41]
 [42 43]]]
```

```
[[[ 0  1]
   [10 11]
   [20 21]
   [30 31]
   [40 41]]
```

```
[[ 2  3]
 [12 13]
 [22 23]
 [32 33]
 [42 43]]]
```

## np.transpose()

It returns an array with axes transposed.

```
In [90]: # with a 2-d array
a = np.linspace(0,250,20).reshape(2,10)
print(a)

print(np.transpose(a))

[[ 0.          13.15789474  26.31578947  39.47368421  52.63157895
   65.78947368  78.94736842  92.10526316 105.26315789 118.42105263]
 [131.57894737 144.73684211 157.89473684 171.05263158 184.21052632
  197.36842105 210.52631579 223.68421053 236.84210526 250.          ]]
```

```
In [91]: # example with more than 2 dimensions, without specifying axis:
a = np.ones((2, 3, 4, 5))
print(np.transpose(a).shape)

(5, 4, 3, 2)
```

```
In [93]: # Example 2-d, specifying axis
a = np.ones((1, 2, 3))
print(np.transpose(a, (1, 0, 2)).shape) #we are indicating that we want the first :

(2, 1, 3)
```

## Stacking together different arrays

There are different ways. To do it along different axis we can use `np.vstack()` or `np.hstack()` when dealing with 2-dimensional arrays.

```
In [100... a = np.empty([2,2])
b = np.zeros([2,2])

print(np.vstack((a,b)))
print(np.hstack((a,b)))

[[2.12199579e-314  4.67296746e-307]
 [5.98807563e-321  3.79442416e-321]
 [0.00000000e+000  0.00000000e+000]
 [0.00000000e+000  0.00000000e+000]]
[[2.12199579e-314  4.67296746e-307  0.00000000e+000  0.00000000e+000]
 [5.98807563e-321  3.79442416e-321  0.00000000e+000  0.00000000e+000]]
```

We can also stack one dimensional arrays as columns in two dimensional arrays using `np.column_stack()` or the row equivalent `np.row_stack()`.

```
In [101... a = np.empty([2,2])
b = np.zeros([2,1])
print(np.column_stack((a, b)))

[[2.12199579e-314  4.67296746e-307  0.00000000e+000]
 [5.98807563e-321  3.79442416e-321  0.00000000e+000]]
```

## Splitting Arrays

Finally, we can also split one array into multiples arrays using `np.hsplit()` or `np.vsplit()`.

```
In [104... a = np.zeros([4,3])
print(np.hsplit(a, 3))  # split a into three arrays

b = np.zeros([5,7])
print(np.vsplit(b,5))  # split into five arrays

[array([[0.],
        [0.],
        [0.],
        [0.]]) array([[0.],
        [0.],
        [0.],
        [0.]]) array([[0.],
        [0.],
        [0.],
        [0.]])
 [array([[0., 0., 0., 0., 0., 0., 0.]]) array([[0., 0., 0., 0., 0., 0., 0.]]) array([[0., 0., 0., 0., 0., 0., 0.]]) array([[0., 0., 0., 0., 0., 0., 0.]]) array([[0., 0., 0., 0., 0., 0., 0.]])]
```

## NumPy Functions

### Statistics

## np.mean()

It allows to compute the arithmetic mean along the specified axis.

```
In [111... a = np.column_stack((np.zeros([4,2]),np.ones([4,1])))
print(a)

print(np.mean(a,axis=0))    # Prints the mean for each column
print(np.mean(a,axis=1))    # prints the mean of each row
print(np.mean(a))           # prints the mean of all the values of the array

[[0. 0. 1.]
 [0. 0. 1.]
 [0. 0. 1.]
 [0. 0. 1.]]
[0. 0. 1.]
[0.33333333 0.33333333 0.33333333 0.33333333]
0.3333333333333333
```

## np.std()

Like `np.mean()` it allows to compute the standard deviation along the specified axis.

```
In [112... print(np.std(a,axis=0))    # there is no deviation
print(np.std(a))

[0. 0. 0.]
0.4714045207910317
```

## np.var()

It allows to compute the variance along the specified axis.

```
In [113... print(np.var(a,axis=0))

print(np.var(a))

[0. 0. 0.]
0.22222222222222224
```

## np.average()

It allows to compute the weighted average along the specified axis.

```
In [114... data = np.arange(6).reshape((3, 2))

np.average(data, axis=1, weights=[1/4, 3/4])    # average across columns

Out[114]: array([0.75, 2.75, 4.75])
```

**TIP!** Sometimes we might have missing observations. With NumPy we can code those as Not a Number using the **np.nan** constant

## np.nanmean()

It allows to compute the arithmetic mean along the specified axis, ignoring Nans.

```
In [119... a = np.array([[1, np.nan], [3, 4]])
print(a)

print(np.mean(a))      # this will be incorrect
print(np.nanmean(a))   # right way to deal with nans

[[ 1. nan]
 [ 3.  4.]]
nan
2.6666666666666665
```

Similar to `np.nanmean()` we could also use `np.nanstd()` and `np.nanvar()` for the standard deviation and variance counterparts.

## Arithmetic Operations

### np.sum()

It allows to sum array elements over a given axis.

```
In [122... data = np.arange(6).reshape((3, 2))
print(np.sum(data,axis=0))  # The function way
print(data.sum(axis=0))    # The method way

[6 9]
[6 9]
```

### np.cumsum()

Returns the cumulative sum of elements along a given axis.

```
In [124... print(data)
print(np.cumsum(data,axis=0))

[[0 1]
 [2 3]
 [4 5]]
[[0 1]
 [2 4]
 [6 9]]
```

### np.min()

It returns the minimum element along a given axis.

```
In [125... print(np.min(data,axis=0))
print(np.min(data))

[0 1]
0
```

The maximum counterpart is `np.max()`.

## np.sqrt()

It returns the non-negative square-root of an array, element-wise.

In [126...

```
print(data)
print(np.sqrt(data))
```

```
[[0 1]
 [2 3]
 [4 5]]
[[0. 1. ]
 [1.41421356 1.73205081]
 [2. 2.23606798]]
```

## np.exp()

It returns the exponential of all the elements in the input array

In [127...

```
print(np.exp(data))
```

```
[[ 1. 2.71828183]
 [ 7.3890561 20.08553692]
 [54.59815003 148.4131591 ]]
```

## Copmarisons

### np.any()

Test wether any array element along a given axis evaluates as True.

In [128...

```
np.any([[True, False], [True, True]])
```

Out[128]: True

On its own it is not very informative, unless we combine it with logical operations, to gain information on our array.

In [129...

```
print(np.any(data>1)) # we have some values greater than 1
```

True

In [130...

```
print(np.any(data>1,axis=0)) # we can also specify the axis
```

```
[ True  True]
```

### np.all()

It is similar to `np.any()` but now evaluates if all the elements of the array satisfy the condition.

In [131...

```
a = np.zeros([3,2])
print(np.all(a==0))
```

True

In [132...

```
# suppose we have a one:

a[0,0] = 1
print(np.all(a==0,axis=0))

[False  True]
```

## Other usefull functions

There are some functions that do not belong to any of the previous categories, but still they are useful and worth studying.

### np.sort()

Returns a sorted copy of any array

In [136...

```
a = np.array([1,435,3,45,65,4,7,65,65756,878,2])
print(a)
print(np.sort(a))

[ 1 435  3 45 65  4  7 65 65756 878  2]
[ 1  2  3  4  7 45 65 65 435 878 65756]
```

### np.round()

Rounds an array to the given number of decimals.

In [141...

```
a = np.array([3.4545,4.3435,6.746757,8.67324])
print(a)

print(np.round(a,decimals=2))

[3.4545  4.3435  6.746757 8.67324 ]
[3.45 4.34 6.75 8.67]
```

### np.where()

Allows to locate elements that satisfy a condition.

In [ ]:

```
a = np.arange(10).reshape([2,5])
print(a)

np.where(a>5)
```

What does the output mean? It is returning the indexes of the elements that satisfy the condition. The first index corresponds to the row, and the second to the column.

# NumPy random sampling

NumPy has a module that produces random numbers. The main functions are:

## `np.random.randint()`

It returns random numbers within the specified values.

```
In [147... a = np.random.randint(5,25,(2,3)) # produces a 2x3 array with numbers in the interval [5,25)
print(a)

[[ 5 23  8]
 [ 9  7 21]]
```

## `np.random.rand()`

Returns random numbers for the specified array dimensions.

```
In [149... a = np.random.rand(2,2) # creates a random 2x2 array
print(a)

[[0.09201743 0.79882103]
 [0.6226247  0.94176854]]
```

## `np.random.random_sample()`

Returns random floats on the half-open interval `[0,1)`.

```
In [150... a = np.random.random_sample((2,3))
print(a)

[[0.80948258 0.42618003 0.58292804]
 [0.49953237 0.65621386 0.3936935  ]]
```

## `np.random.seed()`

It allow us to set the random number generator seed to achieve replicability of our results.

```
In [153... np.random.seed(10) # set seed to 10

a = np.random.randint(0,20,(2,3))
print(a)

[[ 9  4 15]
 [ 0 17 16]]
```

## `np.random.uniform()`

It allows to draw from a uniform distribution in the prespecified interval.

```
In [156... a = np.random.uniform(3,7,(2,2)) # a 2x2 array in the interval [3,7] drawn from a uniform distribution
print(a)

[[3.10068691 5.83683204]
 [4.06226451 4.05441138]]
```

## np.random.binomial()

Allows to draw from a binomial distribution. We can choose the amount of draws and the probability of a 1.

```
In [158... a = np.random.binomial(5,0.3,size=20) # a 20x1 array of 5 repetated draws with su
print(a)

[2 0 2 1 2 1 1 0 3 1 0 2 0 2 2 1 2 3 4 1]
```

## np.random.normal()

Allows to draw from a normal distribution with a given mean and standard deviation.

```
In [159... mean = 5
std = 2
size = (4,5)
a = np.random.normal(mean,std,size)
print(a)

[[4.24490727 6.39126232 7.0496893  8.41551089 1.06576555]
 [3.25808667 2.78634904 5.55024732 4.27108771 2.66969306]
 [2.31386023 7.90678572 5.0023389  4.97615923 6.43176629]
 [2.62846291 3.72734419 3.77837231 2.36121721 3.87654317]]
```

# Save and Load Arrays

We have already seen most of the functions that allow us to work with arrays. But how can we save an array to our computer?

## np.savetxt()

It allows to save an array in '.txt' format. It has an important limitation, arrays can only have 1 or 2 dimensions.

```
In [161... # save the previous array
np.savetxt('array.txt',a) # we are saving with the name array.txt
```

## np.loadtxt()

To load an array already saved, we can just use `np.loadtxt()`.

```
In [162... newarray = np.loadtxt('array.txt')
print(newarray)

[[4.24490727 6.39126232 7.0496893  8.41551089 1.06576555]
 [3.25808667 2.78634904 5.55024732 4.27108771 2.66969306]
 [2.31386023 7.90678572 5.0023389  4.97615923 6.43176629]
 [2.62846291 3.72734419 3.77837231 2.36121721 3.87654317]]
```

# Histograms



NumPy also allow us to prepare our data to future visualizations. We can for example compute histograms with `np.histogram()` which allows to compute the histogram of a dataset.

```
In [163... a = np.histogram([1, 2, 1], bins=[0, 1, 2, 3]) # we input an array and the bins
print(a)
```

```
(array([0, 2, 1], dtype=int64), array([0, 1, 2, 3]))
```

The function returns the values of the histogram and the bin edges. We can also obtain the histogram normalized as a density.

```
In [164... a = np.histogram([1, 2, 1], bins=[0, 1, 2, 3], density=True)
print(a)
```

```
(array([0.        , 0.66666667, 0.33333333]), array([0, 1, 2, 3]))
```

```
In [ ]:
```